



NCCL

NCCL

- 特点：
 - multi-GPU Collective communication primitives
 - topology-aware: 会自动判断拓扑结构

Overview of NCCL

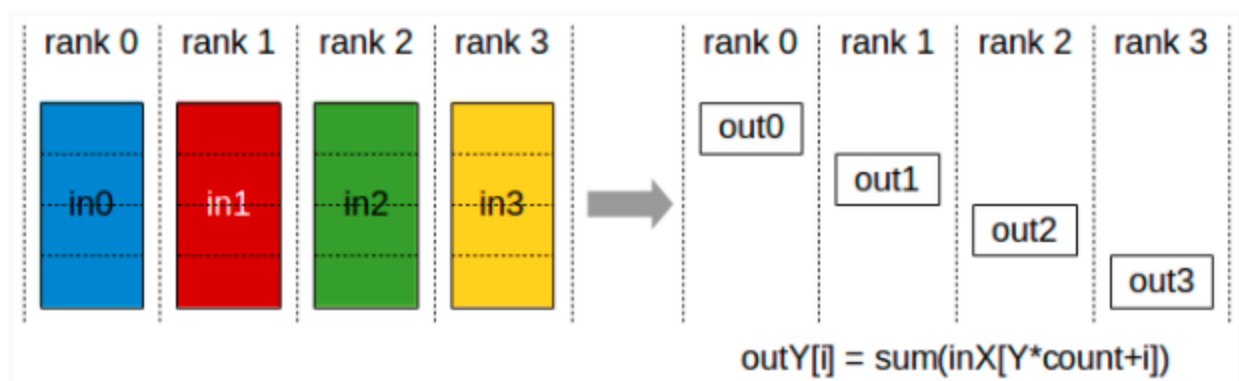
- NCCL provides the following collective communication primitives:
 - AllReduce
 - Broadcast
 - Reduce
 - AllGather
 - ReduceScatter
- 好像会比较好学：文档里的原话：Anyone familiar with MPI with thus find NCCL's API very natural to use.

Creating a Communicator

- First call `ncclGetUniqueId` : create a unique object which will be used by all processes and threads to synchronize and understand they are part of the same communicator.
- Then `ncclCommInitRank` and `ncclCommInitAll` will: create communicator objects, each communicator object being associated to a fixed rank.
 - `ncclCommInitRank()` Creates a new communicator, each rank is associated to a CUDA device.
- `ncclCommDestroy()` : destroy a Communicator

Collective Operations

- AllReduce: `ncclAllReduce()`
- Broadcast: `ncclBroadcast()`
- Reduce: `ncclReduce()`
- AllGather: `ncclAllGather()`
- ReduceScatter: `ReduceScatter()`



Group Calls

- 把很多的Calls合并为一个（提高性能，防止死锁）
- `ncclGroupStart`, `ncclGroupEnd`

- 提高性能的一个点是：Aggregated Operations, this is useful for reducing the launch overhead, as it only occurs once for multiple operations.

Point-to-point communication

`ncclSend()` , `ncclRecv()`

In-place Operations

For `ncclBroadcast`, `ncclReduce` and `ncclAllreduce` functions, this means that passing `sendBuff == recvBuff` will perform in place operations, storing final results at the same place as initial data was read from.

其他的一些API

- `ncclCommAbort()` : Frees resources that are allocated to a communicator object *comm*. Will abort any uncompleted operations before destroying the communicator.
- `ncclCommCount()` : Returns in count the number of ranks in the NCCL communicator *comm*.
- `ncclCommCuDevice()` : Returns in *device* the CUDA device associated with the NCCL communicator *comm*
- `ncclCommUserRank()` : Returns in *rank* the rank of the NCCL communicator *comm*
- 一些Communication Functions: <https://docs.nvidia.com/deeplearning/nccl/user-guide/docs/api/colls.html>

NCCL中的一些Types

<https://docs.nvidia.com/deeplearning/nccl/user-guide/docs/api/types.html>
